

Fast* Implementations of Finite Fields and Elliptic Curves in Lean 4

Matej Penciak

Lurk Lab

March 23, 2023

Slides available at:

TODO: INCLUDE LINK



- ① Motivation: what we're building at Lurk Lab
- ② A bit about zero knowledge cryptography
- ③ Optimized FF arithmetic
 - Add chains for fast exponentiation
 - Batched inversion
 - Montgomery form
- ④ Optimized Elliptic Curve arithmetic
 - Point representation
 - Efficient endomorphisms
 - Multi-scalar multiplication
- ⑤ Brief demo!
- ⑥ Lessons learned and future outlook

Motivation

Who am I? Software engineer and aspiring cryptographer at Lurk Lab (Formerly Yatima).

What is Lurk Lab? We're building the Lurk programming language, and an ecosystem of tools around it. Examples: The Yatima compiler and content addresser, a Lean typechecker, and the cryptography to generate "Lurk proofs"

What is Lurk? Lurk is a turing complete programming language in the style of Lisp which generates zero knowledge proofs of execution.

If I run a Lurk program f , I can prove things like:

- I have run the computation of f applied to u and gotten v
- I know an input u to f such that $f(u) = v$ without revealing u
- I can commit ahead of time to a program f and run it on any inputs $\{u\}$ you provide with proofs of execution.

A little about zero knowledge proofs

Modern perspective on zero knowledge proofs: they are built out of two components

- An interactive proof scheme (the “information-theoretic” component)
- A commitment scheme (the “cryptographic” component)

An IP scheme is a protocol between two parties P (prover) and V (verifier) where (the potentially untrustworthy) P can provide a convincing argument to V that they know some fact/have performed some computation.

An example IP

P and V agree to two graphs G_1 and G_2 , and P wants to prove to V that they know an isomorphism $\phi : G_1 \rightarrow G_2$ without revealing the isomorphism.

- 1 P chooses a random permutation σ of G_1 , and sends $H = \sigma(G_1)$ to V
- 2 V chooses randomly between $\{1, 2\}$ and sends the choice i to P
- 3 P responds with an isomorphism $\phi_i : G_i \rightarrow H$
- 4 V checks that in fact $\phi_i(G_i) \cong H$
- 5 Repeat until V is sufficiently convinced

This protocol is **complete**, **knowledge sound**, and **zero knowledge**

It is not: **non-interactive**, or **succinct**

The goal is to build **Zero Knowledge Succinct Arguments of Knowledge** (ZK SNARKs) in Lean.

The basic building blocks for cryptography

Add some cryptography to the IPs:

Work over finite fields, and encode our computations in so-called “arithmetic circuits”.

The level of security necessary is measured in the number of “bits” of security. To prevent a dishonest prover from forging a fraudulent proof, a good estimate is ≈ 128 bits of security.

Example 1: The fastest (publicly) known factorization/discrete logarithm algorithm is the general number field sieve method with complexity

$$\exp \left(\left(\sqrt[3]{\frac{64}{9}} + o(1) \right) (\log n)^{\frac{1}{3}} (\log \log n)^{\frac{2}{3}} \right)$$

To achieve ≈ 128 bits of security we need to work over finite fields of size $\approx 2^{2048}$

Example 2: Let E be an elliptic curve with a cyclic subgroup of points of size N .

The best (publicly) known discrete logarithm algorithm (given $P = a \cdot G$, find $a \in \mathbb{N}$) is the Pollard- ρ method which has complexity $O(\sqrt{N})$.

To achieve ≈ 128 bits of security we only need a curve with $\approx 2^{256}$ number of points.

Thanks to Hasse's bounds, it's "easy" to find elliptic curves with enough points, if E is an elliptic curve defined over a finite field with q elements, then the number of points N is bounded by:

$$|N - (q + 1)| \leq 2\sqrt{q}$$

When we discuss curves, we will see some magical curves with extra nice properties

Big numbers in Lean

Most computers don't natively support arithmetic for numbers larger than $2^{64} - 1$, stored in the Lean datatype `UInt64`.

Lean 4's `Nat` and `Int` types are built on multiprecision arithmetic libraries written in C(++), and linked in via the `@[extern "..."]` attribute

```
@[extern "lean_nat_add"]
protected def Nat.add : (@& Nat) → (@& Nat) → Nat
  | a, Nat.zero    => a
  | a, Nat.succ b => Nat.succ (Nat.add a b)
```

which calls the C function defined in `lean.h`

```
static inline lean_obj_res
  lean_nat_add(b_lean_obj_arg a1, b_lean_obj_arg a2) {
  . . .
}
```


Multiprecision arithmetic libraries

Lean uses the **GNU MultiPrecision** library on Linux based systems. Mac and Windows get a home-brewed version written in the Lean 4 C++ runtime code.

Both implementations represent Nats as arrays of UInt64s, and define arithmetic in base 2^{64} . If n is the number of “limbs” (`uint64_ts` in C) then the efficiency of some of the most commonly employed algorithms are:

- 1 Schoolbook multiplication: $O(n^2)$
- 2 Karatsuba multiplication: $O(n^{1.58})$
- 3 Strassen FFT algorithm: $O(n \log n \log \log n)$

Because our goal is to work in a fixed-precision finite field, we tried replacing the natural choice of `Nat` with a fixed precision arithmetic based off Lean's `ByteArray` type.

This turned out to be $\approx 100\times$ slower than `Nat`, so we quickly dropped that plan.

A comment on benchmarking and testing

Testing and benchmarking is very important for non-formalized systems. For testing we have built a testing framework called LSpec which is expressive and easy to use.

```
#lspec check "add_comm"  
  (forall n m : Nat, n + m = m + n)  
#lspec check "not true"  
  (forall n m : Nat, n^2 + m^2 >= n^3 + m)
```

which returns

```
? add_comm  
x not true  
=====  
Found problems!  
n := 2  
m := 0  
issue: 8 <= 4 does not hold  
(1 shrinks)
```

A comment on benchmarking and testing

For benchmarking we have a rudimentary suite of tools for running benchmarks:

```
import YatimaStdLib.Benchmark
import YatimaStdLib.AddChain

open Benchmark Exp

def f1 : Nat → Nat := (37 ^ .)
def f2 : Nat → Nat := fastExp 37
def addChainBench : Comparison f1 f2 where
  inputs := #[10000000, 15000000, ...]

def main (args : List String) : IO UInt32 :=
  benchMain args addChainBench.benchmark
```

Benchmarking continued...


This can be run as `Benchmarks-AddChain --num 50 --log out.bench` and times the benchmarks and writes to the file:

```
10000000: f:213993503 vs g:236587657
15000000: f:315046597 vs g:335218176
20000000: f:457342988 vs g:537591730
25000000: f:598272104 vs g:681122849
30000000: f:670924537 vs g:743308727
35000000: f:786821581 vs g:940891049
40000000: f:950924114 vs g:1103080520
45000000: f:1108323608 vs g:1213150412
50000000: f:1242277288 vs g:1288141664
```

The arithmetic typeclasses we have (so far):

- `Ring`: Basic arithmetic: $+$, $*$, and so on
- `Field`: `Ring` together with inversion
- `PrimeField`: Arithmetic, plus pre-computations and methods to improve efficiency
- `NewField`: Automatically generated instance of a `PrimeField` with special methods related to Montgomery form
- `GaloisField`: Fields obtained as extension fields of a base field

Field typeclasses



(thanks ChatGPT)

Example classes

```
class Ring (R : Type _) extends Add R, Mul R, Sub R,  
    HPow R Nat R, BEq R, Coe Nat R where
```

```
  zero : R
```

```
  one : R
```

```
class PrimeField (K : Type _) extends Field K where
```

```
  char : Nat -- characteristic
```

```
  sqrt : K → Option K -- Efficient `sqrt` implementation
```

```
  content : Nat -- char.log2
```

```
  twoAdicity : Nat x Nat -- `(s, r)` where  $p-1 = 2^s * r$ 
```

```
  legAC : Array ChainStep -- Pre-computed addchains to
```

```
  frobAC : Array ChainStep -- calculate  $p$  and  $(p-1)/2$ 
```

```
  fromNat : Nat → K -- to and from `Nat` methods
```

```
  natRepr : K → Nat
```

```
  batchedExp : K → Array Nat → Array K
```

```
  batchedInv : Array K → Array K
```

AddChains, batched arithmetic, and Montgomery form

AddChains: Exponentiation is a commonly used operation, but can be very slow.

Naive implementation: $a^N = a \cdot a^{N-1}$ is $O(N)$ (too slow)

Square and multiply method requires $\approx \frac{3}{2} \log N$ operations ($\log N$ squarings plus the average Hamming weight of N 's binary expansion) (WARNING: susceptible to side channel attacks)

For common exponents (p and $\frac{p-1}{2}$): pre-compute a minimal length “addition chain”: $[n_1, n_2, n_3, \dots, n_r = N]$ where each $n_i = n_k + n_l$ for $k, l < i$.

Example: Calculating the Frobenius modulo the prime

0x40000000000000000000000000000000224698fc094cf91b992d30ed00000001
381 field ops using square and multiply, and 337 field ops using a minimal addition chain.

Batched inversion

Batched inversion:

Inversion is another costly operation (inversion in $\mathbb{Z}/p\mathbb{Z}$ is $O(\log p)$ using the extended Euclidean algorithm)

To invert $[n_1, n_2, n_3, \dots, n_r]$, calculate $N_j = \prod_{i=1}^j n_i$ for $j = 1, \dots, r$ caching the results. Invert N_r . Recover all of the n_i^{-1} by multiplying with the cached results.

Example: $[2, 3, 5]$.

- 1 Calculate $N_1 = 2$, $N_2 = 6$, $N_3 = 30$.
- 2 Invert N_3 : $N_3^{-1} = 30^{-1}$, let $t = N_3^{-1}$
- 3 Multiply $t \cdot N_2 = 30^{-1} \cdot 6 = 5^{-1}$, and $t = t \cdot n_3 = 6^{-1}$
- 4 Proceed as before: $t \cdot N_1 = 6^{-1} \cdot 2 = 3^{-1}$, and $t = t \cdot n_2 = 2^{-1}$

This method replaces r field inversions with 1 field inversion and $O(3r)$ field multiplications.

Montgomery form

Problem: How should we calculate $n \cdot m$ in \mathbb{F}_p when we use Nat representatives?

Naive algorithm: Calculate $a \cdot b \bmod p$, but this requires division by p for every multiplication.

Intuition behind the solution: Replace reduction modulo p with reduction modulo R for some other R .

Let $R > p$ be a number co-prime to p for which division and reduction modulo R is efficient.

In practice: Choose R to be a power of 2 greater than p .

Montgomery reduction is an efficient algorithm that calculates $nR^{-1} \bmod p$ using only arithmetic modulo R .

Represent n in "Montgomery form" $[n]_w = n \cdot R \bmod p$.

Addition: $[n]_w + [m]_w = (n + m) \cdot R \bmod p = [n + m]_w$

Multiplication: $[n]_w \cdot [m]_w = (n \cdot m)R^2 \bmod p = [n \cdot m]_w R \bmod p$

NewField implementation

We implement a `macro_rule` that expands the following declaration

```
new_field TestField with
  prime: 2011
  generator: 7
  root_of_unity: 2010 -- optional
```

into a type `TestField` with pre-computations and arithmetic, instances, and methods implemented.

Not advised to be used yet: Missing some key GMP methods to make Montgomery reduction more efficient (`mpz_tdiv_r_2exp`).

Part of Montgomery reduction of $a \in \mathbb{N}$ involves calculating $a \bmod R$.

Elliptic Curves

Short term goals have us working with curves defined over fields of large characteristic, so representing a curve is simple:

```
--  
Curves with Weierstrass form satisfying the equation  
`y2 = x3 + a x + b` for a prime field `F` such that  
`char K > 3`  
-/  
structure Curve (F : Type _) [Field F] where  
  a : F  
  b : F
```

Arithmetic on curves

Arithmetic is more interesting:

```
class CurveGroup {F : outParam $ Type _} [Field F] (K : Type _)  
  zero : K  
  inv : K → K  
  add : K → K → K  
  double : K → K
```

Separating the curve from its arithmetic is important to allow for different coordinate systems on the same curve.

Projective coordinates avoid field inversions!

Affine coordinates are more compact on disk!

More exotic coordinate systems have other benefits.

Coordinate breakdown

Doubling		Addition	
Operation	Costs	Operation	Costs
$2\mathcal{P}$	7M + 5S	$\mathcal{J}^m + \mathcal{J}^m$	13M + 6S
$2\mathcal{J}^c$	5M + 6S	$\mathcal{J}^m + \mathcal{J}^c = \mathcal{J}^m$	12M + 5S
$2\mathcal{J}$	4M + 6S	$\mathcal{J} + \mathcal{J}^c = \mathcal{J}^m$	12M + 5S
$2\mathcal{J}^m = \mathcal{J}^c$	4M + 5S	$\mathcal{J} + \mathcal{J}$	12M + 4S
$2\mathcal{J}^m$	4M + 4S	$\mathcal{P} + \mathcal{P}$	12M + 2S
$2\mathcal{A} = \mathcal{J}^c$	3M + 5S	$\mathcal{J}^c + \mathcal{J}^c = \mathcal{J}^m$	11M + 4S
$2\mathcal{J}^m = \mathcal{J}$	3M + 4S	$\mathcal{J}^c + \mathcal{J}^c$	11M + 3S
$2\mathcal{A} = \mathcal{J}^m$	3M + 4S	$\mathcal{J}^c + \mathcal{J} = \mathcal{J}$	11M + 3S
$2\mathcal{A} = \mathcal{J}$	2M + 4S	$\mathcal{J}^c + \mathcal{J}^c = \mathcal{J}$	10M + 2S
—	—	$\mathcal{J} + \mathcal{A} = \mathcal{J}^m$	9M + 5S
—	—	$\mathcal{J}^m + \mathcal{A} = \mathcal{J}^m$	9M + 5S
—	—	$\mathcal{J}^c + \mathcal{A} = \mathcal{J}^m$	8M + 4S
—	—	$\mathcal{J}^c + \mathcal{A} = \mathcal{J}^c$	8M + 3S
—	—	$\mathcal{J} + \mathcal{A} = \mathcal{J}$	8M + 3S
—	—	$\mathcal{J}^m + \mathcal{A} = \mathcal{J}$	8M + 3S
—	—	$\mathcal{A} + \mathcal{A} = \mathcal{J}^m$	5M + 4S
—	—	$\mathcal{A} + \mathcal{A} = \mathcal{J}^c$	5M + 3S
$2\mathcal{A}$	I + 2M + 2S	$\mathcal{A} + \mathcal{A}$	I + 2M + S

Naiive vs. Actual implementation

```
let mut X3 := a * Z3
let mut Y3 := b3 * t2
Y3 := X3 + Y3
X3 := t1 - Y3
Y3 := t1 + Y3
Y3 := X3 * Y3
X3 := t3 * X3
Z3 := b3 * Z3
t2 := a * t2
t3 := t0 - t2
t3 := a * t3
t3 := t3 + Z3
Z3 := t0 + t0
t0 := Z3 + t0
t0 := t0 + t2
t0 := t0 * t3
Y3 := Y3 + t0
t2 := Y * Z
t2 := t2 + t2
t0 := t2 * t3
X3 := X3 - t0
Z3 := t2 * t1
Z3 := Z3 + Z3
Z3 := Z3 + Z3
return (X3, Y3, Z3)
```

Reason: Ensure Lean is evaluating the arithmetic expressions to minimize finite field arithmetic

Scalar multiplication and GLV optimization

Given $n \in \mathbb{N}$ and $P \in C$, want to calculate $n \cdot P$.

If C has prime order, and an efficiently computable endomorphism

$\Phi : C \rightarrow C$ then $\Phi(P) = \Lambda \cdot P$ for some $\Lambda \in \mathbb{N}$

Use this to compute

$$n \cdot P = (n_1 + \Lambda n_2) \cdot P = n_1 \cdot P + n_2 \cdot \Phi(P)$$

where n_1 and n_2 are $O(\sqrt{n})$

This optimization was subject to the patent US7995752B2 “Method for accelerating cryptographic operations on elliptic curves” which expired in 2020

If the curve C is defined as $y^2 = x^3 + b$ then $(x : y : z) \mapsto (\zeta x : y : z)$ for $\zeta \in \mu_3$ is a candidate endomorphism.

Some very magical curves

The curve is $y^2 = x^3 + 5$ over

0x400000000000000000000000000000000224698fc094cf91b992d30ed00000001
(Pallas) and

0x400000000000000000000000000000000224698fc0994a8dd8c46eb2100000001
(Vesta)

Nice properties of the Pasta curves:

- 1 The group of points on one is cyclic of order equal to the characteristic of the other's base field.
- 2 Efficiently computable endomorphism for the GLV optimization.
- 3 Large 2-adicity ($p - 1$ is divisible by 2^{32})
- 4 Both primes are of the form $2^{255} + \epsilon$ which helps with reduction in \mathbb{F}_p
- 5 Both have low-degree isogenies with non-zero j -invariant (this helps with generating random points on the curves).

The story of how the curve was discovered is fascinating.

Multiscalar Multiplication

The multi-exponentiation problem is the following: Given a set of list of integers $[n_1, n_2, \dots, n_r]$ and a list of group elements $[g_1, g_2, \dots, g_r]$ calculate $\prod_i^r g_i^{n_i}$.

Multi-exponentiation in the context of elliptic curve groups asks to calculate $\sum_i^r n_i \cdot G_i$ for points $G_i \in C$.

Multiscalar multiplication is used throughout cryptography (signature schemes, commitment schemes, zero knowledge proof schemes).

In some protocols, $\approx 80\%$ of time in generating zero knowledge proofs is spent calculating multiscalar multiplications.

Typical r varies greatly by application, but can get up to sizes like 2^{28} .

Pippenger's algorithm

Pippenger's algorithm is one of the most efficient optimizations over the naive implementation.

Suppose the elliptic curve group order is approximately on the order of b bits. Pick a “window size” c ($\log r$ turns out to be optimal) so that $b = k \cdot c$.

- 1 Decompose each $n_i = n_{i,0} + n_{i,1} \cdot 2^c + n_{i,2} \cdot 2^{2c} + \dots + n_{i,k} \cdot 2^{kc}$
- 2 Split up the large b -bit MSM into k separate c -bit MSMs:

$$T_i = \sum_j^r n_{j,i} G_i$$

- 3 Solve each c -bit MSM and recombine as $M = \sum_i^k 2^i \cdot T_i$.
(perfect opportunity for parallelization)

Pippenger algorithm

We have reduced to the case of an MSM $T = \sum_i^r n_i \cdot G_i$ where each n_i is at most c -bits.

- 1 Keep track of $2^c - 1$ “buckets”, and place each G_i into the n_i bucket.
- 2 Sum up all the sets of points to get 2^c bucket sums S_j for $j = 0, \dots, 2^c - 1$
- 3 $T = S_{2^c-1} + (S_{2^c-1} + S_{2^c-2}) + \dots + (S_{2^c-1} + S_{2^c-2} + \dots + S_0)$

The Lean implementation provides an $\approx 10\times$ speedup over the naive implementation.

Further optimizations exist when each G_i is a fixed shape relative to a generator for the elliptic curve group.

A common case is when $G_i = \tau^i \cdot G$ for $i = 0, \dots, r$ and G is a fixed generator for the elliptic curve group.

A small Demo!

Lessons learned

- 1 Squeezing performance out of Lean can feel like squeezing water out of stone sometimes.
- 2 Lean 4 does not have “zero cost abstractions”.
- 3 Lean 4 is early in its life, so there are few to no guides written down for writing efficient code.
- 4 There are some significant performance pits one can fall in if one isn't careful with the reference counting.
- 5 Lean 4 is incredibly expressive, it shines in terms of the ease by which one can prototype complex structures and algorithms.

Future goals

- ➊ (short term) Optimize, optimize, optimize! (deeper GMP integration, more curve forms, ...)
- ➋ (short-medium term) Expand FFaCiL to include more elliptic curve cryptography (isogenies!)
- ➌ (medium term) Begin exploring the formalization of provable security (formalizing adversarial games, and complexity)
- ➍ (long term) Expand the Yatima Standard Library with more number theoretic primitives.
- ➎ (Long term) Begin the process of formally verifying these algorithms!
- ➏ (Epochs from now) Combine all of these efforts together into a `cryptolib` with formally verified cryptographic algorithms.

Thank you for your time!

Thanks!

TODO